

Introduction to Pathfinding I


BFS and DFS

Ryan Toner

Basic Graph Pathfinding - Learning Objective

- Graphs are a data structure that consist of the following:
 - We define the graph G in the following way
 - $V(G) = \{v_1, v_2 \dots v_k\}$, a set of k vertices, or nodes
 - $E(G) = \{e_1, e_2 \dots e_j\}$, a set of j edges, where e_i defines a tuple $e_i = (v_a, v_b)$, the edge between two vertices V_a and V_b
- A path in the graph may be thought of simply as a route existing between two nodes by following all legal edge connections
 - They should not repeat vertices or edges (no cycles)
 - They should follow an ordered direction from start \rightarrow finish
- The objective of this lecture is to introduce the most fundamental graph searches - DFS and BFS.

Example Graph



- This is a directed graph.
- $V(G) = \{A, B, C, D, E, F, G\}$
- $E(G) = \{(A, D), (B, G), (C, E), (D, B), (D, C), (F, E), (G, A), (G, F)\}$
- Suppose we would like to know if we could find a path given an arbitrary starting and ending node.

Graph Search - DFS


- The first algorithm for graph search we will look at is Depth-First search (DFS).
- This search will take a source node (root) and will find paths to an end node (terminus).
- This search can easily be extended to find all possible paths in the graph starting from the start (simply by not specifying an ending node).

High-Level DFS Algorithm:


1. Push root node to stack
2. Pop the top node off the stack
 - a. If the top node is marked, ignore it and repeat step 2
3. Mark the top node*
4. Push all unmarked neighbors of the top node to the stack
5. Repeat step 2-5, until stack is empty or terminus point is found

*A Note about marking:

- You can represent marking with a hashset for efficiency e.g. marked items are in the set
- Marking the node at step 3 may save time in graphs with self-directed edges, where the edge connects a node to itself.




Graph Search - DFS



- Use DFS with a stack and marked list to find a path from A to F
- The stack will contain tuples of the form (node, path), where node indicates the node we have reached, and path indicates an augmenting path of how we got there.
- Start by placing the starting node on the stack.
 - Push (A, null) onto the stack. The null value indicates that this is the starting point.



Graph Search - DFS





- Pop (A, null) from the top of the stack.
- Mark A (blue outline) to ensure it is only processed once.
- Next, process the unmarked neighbors of A.
- The neighbors of A are:
 - **all** edges **e** such that **e=(A,_)**, where **_** is any connected node.
- In this case, we only find D.
- Push (D, {A}) to the stack.





Graph Search - DFS

1. Pop (D, {A}) from the top of the stack.
2. Mark D
3. Process the unmarked neighbors of D:
 - a. Push (B, {A,D}) to the stack.
 - b. Push (C, {A,D}) to the stack.




Graph Search - DFS

1. Pop $(C, \{A,D\})$ from the top of the stack.
2. Mark C
3. Process the unmarked neighbors of C:
 - a. Push $(E, \{A,D,C\})$ to the stack.




Graph Search - DFS

1. Pop (E, {A,D,C}) from the top of the stack.
2. Mark E
3. Process the unmarked neighbors of E:
 - a. E has no neighbors



Graph Search - DFS


1. Pop (B, {A,D}) from the top of the stack.
2. Mark B
3. Process the unmarked neighbors of B:
 - a. Push (G, {A,D,B}) to the stack.



Stack

Graph Search - DFS


1. Pop (G, {A,D,B}) from the top of the stack.
2. Mark G
3. Process the unmarked neighbors of B:
 - a. Push (F, {A,D,B,G}) to the stack.
 - b. A is already marked, so do not push.



Stack


Graph Search - DFS

1. Pop (F, {A,D,B,G}) from the top of the stack.
2. Mark F
3. F is the terminus node, so terminate the algorithm here.



Stack

Graph Search - DFS



1. Our algorithm returns $(F, \{A, D, B, G\})$.
2. We can augment F to the path for the final result: $\{A, D, B, G, F\}$ as our full path from A to F .
3. Another way this algorithm could work would be by simply storing the parents of each node in the stack.
4. An abstract data type is needed to store a node and its parent




Stack

Simple Abstract Data Type (NodePath)


- This ADT will have two fields:
 1. Node n;
 2. NodePath parent;
- The NodePath will simply store the parent that got to Node n
- Suppose we ran our DFS using this method instead of (Node n, Path {})

Graph Search - DFS with NodePath ADT




- Our algorithm would actually produce a linked-list!
- We would see something like this:
- $F \rightarrow G \rightarrow B \rightarrow D \rightarrow A \rightarrow \text{Null}$
- Notice that the arrow direction has flipped!
- Consider why, starting from the beginning:
- The first element, $p_1 = (A, \text{null})$ is still the same since A has no parent
- Then $p_2 = (D, p_1)$ is the second
- Then $p_3 = (C, p_2)$ is the third
- So, by the time we get to F, we have $p_n = (F, p_{n-1})$. However, we must traverse the linked list backwards now to determine the path. This technique is called “**backtracking**” and is critically important in Graph Theory.


Graph Search - DFS with NodePath ADT




- We need to get the path in order from A→F
- First, add the entire linked-list path into a stack
- $P = (F, p_{n-1})$
- while($P \neq \text{null}$):
 - `stack1.push(P)`
 - $P = P.\text{parent}$
- For simplicity, I will label all parents in numerical order (though this is not the actual order of traversal)



Graph Search - DFS with NodePath ADT



- Now, we can read off the stack in-order and get the actual path from A to F
 - `Path = list()`
 - `while(stack is not empty):`
 - `Path.append(stack.pop().n)`
- ⇒
- `Path = {A,D,B,G,F}`
 - We have our path in-order again!



Overview - DFS

- The name of this pathfinding algorithm is Depth-First Search (DFS)
- Notice the way it looked through the first path (preferring depth) before going on to to the second path.
- DFS is good at making long gains quickly, but it may waste time if it checks paths in a inefficient ordering.
- The usage of the stack is necessary for DFS to happen.
- Using a “marking” system saves computation speed by making sure we do not recheck already processed or currently processing nodes. This can be accomplished easily with a HashSet.

Graph Search - BFS


- The second algorithm for graph search we will look at is Breadth-First search (BFS).
- This search will take a source node (root) and will find paths to an end node (terminus).
- This search can easily be extended to find all possible paths in the graph starting from the start (simply by not specifying an ending node).

High-Level BFS Algorithm:

1. Enqueue root node to queue
2. Dequeue the top node from the queue
 - a. If the top node is marked, ignore it and repeat step 2
3. Mark the top node*
4. Enqueue all unmarked neighbors of the top node to the queue
5. Repeat step 2-5, until queue is empty or terminus point is found


*The same note about marking still applies:

- You can represent marking with a hashset for efficiency e.g. marked items are in the set
- Marking the node at step 3 may save time in graphs with self-directed edges, where the edge connects a node to itself.



Graph Search - BFS

- Let's try this again, except simply swap out the Stack for a Queue data structure. Remember that queues are first-in, first-out (FIFO).
- This will process nodes in the order they appear.
- Enqueue (A, null)




(A, null)

Queue

Graph Search - BFS

1. Dequeue (A, null) from the front of the queue.
2. Mark A
3. Process the unmarked neighbors of A:
 - a. Enqueue (D, {A}).




(D, {A})

Queue

Graph Search - BFS

1. Dequeue (D, {A}) from the front of the queue.
2. Mark D
3. Process the unmarked neighbors of D:
 - a. Enqueue (C, {A,D}).
 - b. Enqueue (B, {A,D}).




(C, {A,D})
(B, {A,D})

Queue

Graph Search - BFS

1. Dequeue (C, {A,D}) from the front of the queue.
2. Mark C
3. Process the unmarked neighbors of C:
 - a. Enqueue (E, {A,D,C}).




(B, {A,D})
(E, {A,D,C})

Queue

Graph Search - BFS

1. Dequeue (B, {A,D}) from the front of the queue.
2. Mark B
3. Process the unmarked neighbors of B:
 - a. Enqueue (G, {A,D,B}).




(E, {A,D,C})
(G, {A,D,B})

Queue

Graph Search - BFS

1. Dequeue (E, {A,D,C}) from the front of the queue.
2. Mark E
3. Process the unmarked neighbors of E:
 - a. E has no neighbors




(G, {A,D,B})

Queue

Graph Search - BFS

1. Dequeue $(G, \{A,D,B\})$ from the front of the queue.
2. Mark G
3. Process the unmarked neighbors of G:
 - a. Enqueue $(F, \{A,D,B,G\})$.
 - b. A is already marked.




$(F, \{A,D,B,G\})$

Queue

Graph Search - BFS


1. Dequeue (F, {A,D,B,G}) from the front of the queue.
2. Mark F
3. F is the terminus node, so terminate the algorithm here.



Queue

Graph Search - BFS

- Our algorithm returns $(F, \{A,D,B,G\})$.
- We can augment F to the path for the final result: $\{A,D,B,G,F\}$ as our full path from A to F.
- Like DFS, we could also use the NodePath ADT approach for BFS and use backtracking to get the ordered solution path.



Queue

Overview - BFS

- BFS may be inefficient when there are many neighbors that require checking before progress further away may be made.
- BFS tends to be the basis for more advanced pathfinding algorithms, but it will use a weight-based approach for selecting neighbors connected by edges with the least-weight.
 - This would use a priority-queue data structure, which returns neighbors in order of highest priority.

BFS vs DFS

- In terms of implementation, the only major difference between BFS and DFS is Queue and Stack usage respectively.
- They are functionally equivalent, but one may outperform the other depending on the graph construction.
- Both may use an augmenting path (list), or NodePath (linked-list) approach to constructing the final-path once the algorithm terminates.
- Either may be used for finding all possible paths in a network from a single-source/root, simply by not specifying a terminal node.
 - The respective algorithm will terminate once the stack or queue is empty.

Pathfinding and Mazes

- DFS tries to make gains to the end as fast as possible
- Humans tend to scan mazes by trying to follow single paths – just like DFS!
- Think about it:
 - You will scan a path until you hit a wall, and then trace back to the next possible path.
 - This is much easier than thinking about solving a maze with BFS, as it is too much information to consider several paths branching out at once!
- Problems that involve lengthy paths are much better with DFS.
 - BFS will take a long time, having to explore all paths node-by-node
 - Try not to use a BFS with a maze! You are better off using informed search algorithms you will learn later.
- However, BFS will find shorter path lengths than DFS.

Conclusions

- BFS and DFS are not too difficult to implement, and a simple change in data structure from Queue to Stack, or vice-versa, will flip your search technique.
- Although they are functionally equivalent, it is sometimes more appropriate to use one over the other
- BFS and DFS are “uninformed” searches. This is because they explore all neighbor nodes without any notion of which is better.